
Parameter Scope Matters: MNE-L2 for Reliable Low-Latency ANN-to-SNN Conversion

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Low-latency artificial-to-spiking neural network (ANN-to-SNN) conversion is at-
2 tive for event-driven edge inference, but robustness depends strongly on how
3 the source ANN is regularized. We identify a parameter-scope failure: applying
4 all-parameter L2 to batch-normalization (BN) affine parameters can make a converted
5 network substantially more sensitive to noise. To explain this behavior,
6 we derive a local boundary-crossing analysis in which representation changes are
7 governed by effective perturbation scale relative to the quantization margin. Guided
8 by this analysis, we propose margin-normalized effective L2 (MNE-L2), which
9 penalizes BN-folded weight amplification relative to that margin while stopping the
10 regularization gradient through BN affine parameters and IF thresholds. Across five
11 CIFAR training seeds, MNE-L2 preserves substantially more accuracy under post-
12 IF noise while reducing clean activity relative to all-parameter L2. On CIFAR-10
13 at $T = 16$, its firing density and estimated arithmetic energy are 0.083 and 0.472
14 mJ, versus 0.105 and 0.585 mJ for all-parameter L2; at $\sigma = 5$, accuracy is 72.06%
15 versus 10.00%. The results expose an accuracy–reliability–activity trade-off rather
16 than universal or hardware-measured dominance.

17 1 Introduction

18 Resource-constrained inference is limited by dense computation, memory movement, and response
19 latency. Spiking neural networks (SNNs) offer an event-driven alternative whose sparse activity can be
20 exploited on neuromorphic hardware [1]. ANN-to-SNN conversion is particularly practical because
21 it retains conventional ANN optimization and maps trained activations to spike-rate representations
22 [2, 3]. Its central difficulty is finite-time fidelity: at small simulation time T , spike counts may not
23 reproduce the source activation, while increasing T directly increases latency and event cost.

24 Quantization-based conversion addresses this accuracy–latency tension by shaping the source ANN
25 activation before conversion [4, 5]. In particular, QCFs uses a quantized clip-floor-shift activation
26 aligned with finite-time firing rates. Yet its discrete representations remain vulnerable when conver-
27 sion residuals or internal perturbations cross a quantization boundary. How training choices alter
28 these crossings and their downstream amplification remains underexplored.

29 We approach this question through the robustness of conversion-based SNNs: how does the source
30 ANN’s regularization scope affect discrete representations after conversion? In networks with batch
31 normalization (BN) [6], “L2” is incomplete unless its parameter scope is specified. We distinguish
32 all-parameter L2, weights-only L2, and selective penalties on weights plus BN affine terms. Our
33 controlled experiments reveal a scope-dependent failure: penalizing BN gain γ is associated with

an early collapse under internal noise, whereas penalizing BN bias β alone is not. This observation motivates both the stability analysis and the conversion-aware regularizer developed below.

To explain this failure, we analyze a converted QCFS layer under additive perturbations. A discrete code changes only when the perturbation crosses a quantization boundary, so local risk is controlled by effective noise relative to the nearest margin. Because downstream BN-folded weights affect effective noise while IF thresholds set the QCFS margin, this view connects parameter scope to representation stability. Guided by the analysis, we propose margin-normalized effective L2 (MNE-L2), which penalizes BN-folded amplification relative to the following margin without directly updating BN affine parameters or IF thresholds through the regularization term.

Our contributions are:

- We isolate the L2 scope effect with weights-only, weights-plus-BN- β , weights-plus-BN- γ , and weights-plus-both controls. In this single-seed ablation, BN- γ distinguishes the two unstable scope settings from the two stable settings.
- We derive a local boundary-crossing view and layerwise margin-to-noise diagnostics that connect effective amplification to changes in discrete QCFS representations.
- Guided by this analysis, we formulate MNE-L2 and evaluate it against L1/L2 baselines using five independently trained CIFAR-10 and CIFAR-100 VGG13/16/19 checkpoints, a single ImageNet ResNet-18 checkpoint, and full-test-set VGG16 firing-density and arithmetic-energy estimates.

2 Related Work

SNN training and ANN-to-SNN conversion. High-performance SNNs are commonly obtained either by direct training or conversion. Direct approaches optimize temporal spiking dynamics through surrogate gradients [7–9], but require backpropagation through multiple simulation steps. Conversion instead trains a source ANN and maps its activations to firing rates. Early methods use weight normalization, threshold balancing, and activation normalization [10, 2]; later work scales conversion to VGG and residual architectures [3]. Our work follows conversion and targets the reliability of intermediate quantized representations at low T , rather than proposing a new neuron model or direct-training rule.

Quantization-based low-latency conversion. Finite-time mismatch has been reduced through potential initialization [11], parameter calibration [12], residual membrane-potential correction [13], offset-spike calibration [14], probabilistic temporal alignment [15], and quantized ANN training [4, 5]. Recent methods also learn conversion-error compensation or bridge quantized activations and spiking neurons [16–18]. These approaches mainly modify conversion dynamics, calibration, or activation design. MNE-L2 is complementary: the conversion rule is unchanged, while the source training objective penalizes layerwise amplification relative to the quantization margin. The BN-scope ablation further shows that the parameter set receiving a nominally standard L2 penalty can materially change measured post-conversion reliability.

Stability and noise sensitivity. Perturbation stability has been analyzed through Lyapunov bounds for ANNs [19], membrane-potential dynamics in trained SNNs [20], spectral properties of random LIF networks [21], and classifier-level adversarial robustness after conversion [22]. We study a different target: whether a controlled post-input-IF perturbation changes intermediate QCFS codes and how these changes propagate through a converted network. The margin-to-noise and bin-crossing diagnostics therefore localize representation failures rather than certifying worst-case classifier robustness.

3 From Parameter Scope to Margin-Aware Regularization

3.1 Low-latency QCFS conversion

For activation x , threshold λ , and quantization level L , the source ANN uses

$$q_L(x; \lambda) = \frac{\lambda}{L} \left\lfloor L \operatorname{clip}\left(\frac{x}{\lambda}, 0, 1\right) + \frac{1}{2} \right\rfloor. \quad (1)$$

At inference, an integrate-and-fire (IF) layer emits a length- T spike sequence whose rate represents the quantized activation. Unless stated otherwise, we use $L = T = 16$ and a rate-uniform schedule for the VGG experiments. Additional depth-transfer and compact low-latency controls are reported only in the appendix.

3.2 Local noise-to-margin analysis

Let z be the clean output of an IF layer and $\tilde{z} = z + \epsilon$, where $\epsilon \sim \mathcal{N}(0, \sigma^2 I)$. Let $\mathcal{B}_l$ contain the QCFS boundaries of layer l and $d_l(u) = \min_{b \in \mathcal{B}_l} |u - b|$. Under a local additive model $\delta_l \sim \mathcal{N}(0, \sigma_{\text{eff},l}^2)$, changing the quantization bin requires crossing a boundary, so

$$\Pr[q_L(u + \delta_l) \neq q_L(u) \mid u] \leq 2Q\left(\frac{d_l(u)}{\sigma_{\text{eff},l}}\right), \quad (2)$$

where Q is the Gaussian upper-tail function. We estimate $\sigma_{\text{eff},l} = \sqrt{\mathbb{E}[(z_l - \tilde{z}_l)^2]}$, summarize $d_l(u)$ by its median $d_{50,l}$, and report

$$\rho_l = \frac{d_{50,l}}{\sigma_{\text{eff},l} + \varepsilon}, \quad P_{\text{cross},l} = \Pr[q_L(z_l) \neq q_L(\tilde{z}_l)]. \quad (3)$$

Small ρ_l and large $P_{\text{cross},l}$ identify layers where noise is likely to change discrete spike counts. In an idealized independent aggregation model, σ_{eff}^2 is upper-bounded by $O(1/(CT))$, while the nominal QCFS half-cell margin decreases as $1/L$; hence a nominal stability proxy scales as $\sqrt{CT}/L$ (Appendix A). This motivates evaluating precision, width, and timestep jointly rather than treating L as an isolated accuracy knob. It also suggests controlling the BN-folded amplification that contributes to $\sigma_{\text{eff},l}$ relative to the following QCFS margin.

3.3 MNE-L2 derived from the stability analysis

Guided by the preceding analysis, for each convolution or linear layer followed by BN and an IF activation, we fold the BN scale into the matched weight while stopping the regularization gradient through BN and the IF threshold:

$$\widetilde{W}_{l,o} = \text{sg}\left(\frac{\gamma_{l,o}}{\sqrt{v_{l,o} + \epsilon_{\text{BN}}}}\right) W_{l,o}, \quad M_{\text{eff},l} = \frac{1}{C_l} \sum_o \|\widetilde{W}_{l,o}\|_F^2. \quad (4)$$

Margin-normalized effective L2 is

$$\mathcal{R}_{\text{MNE}} = \sum_l \frac{L_l^2 M_{\text{eff},l}}{\text{sg}(\lambda_l)^2 + \varepsilon}, \quad (5)$$

where sg denotes stop-gradient. Layers without a following IF threshold are excluded. The objective uses the current folded scale to weight each matched W_l , but the penalty does not directly shrink BN- γ , BN- β , or λ_l . This distinction motivates the scope ablation: standard all-parameter L2, weights-only L2, and selective L2 on weights plus BN affine terms are different interventions even when their coefficients are numerically similar.

4 Experiments

4.1 Setup

We use CIFAR-10 and CIFAR-100 [23] with QCFS VGG16 and ImageNet [24] with QCFS ResNet-18 [25]. Unless stated otherwise, $L = T = 16$ and Gaussian noise is injected after the input IF. We denote optimizer L2 applied to all trainable parameters as L2-all. L2-wo restricts the same penalty to convolutional and linear weight tensors, excluding biases, BN affine parameters, and trainable IF thresholds; L1-wo uses the same weights-only scope. The parameter-scope ablation uses CIFAR-10 seed 42 and one stored perturbation stream at each $\sigma \in [0, 5]$. It compares weights-only L2, weights plus BN- β , weights plus BN- γ , and weights plus both affine terms. The cross-dataset CIFAR study evaluates five independently trained checkpoints (seeds 40–44) for L2-all, L2-wo, L1-wo, and MNE-L2. Curves show means and sample standard deviations across training seeds.




Figure 1: CIFAR-10 VGG16 L2 parameter-scope ablation at $L = T = 16$. Gaussian noise is injected after the input IF. The training seed is fixed at 42 and each point uses one stored perturbation stream, so no uncertainty band is shown. Regularizing BN- γ , with or without BN- β , causes the early collapse.

118 Appendix C.2 repeats the CIFAR sweep with VGG13 and VGG19 and includes a fixed- α Calibrated
 119 MNE control.

120 The scale-transfer experiment evaluates the same four methods on one ImageNet ResNet-18 check-
 121 point (seed 42) over $\sigma \in [0, 2]$. Thus, the ImageNet result is a single-checkpoint extension, not a
 122 variance-aware benchmark. Hyperparameters are fixed before the noise evaluation, and the reported
 123 configurations are not clean-accuracy matched. The injected perturbation is a synthetic internal-noise
 124 stress test rather than an adversarial input attack [22] or a substitute for device measurement. CIFAR
 125 activity and energy accounting uses all 10,000 test images for each of five seeds at $T \in \{4, 8, 16\}$;
 126 only clean activity is used for deployment energy. Layerwise diagnostics use held-out CIFAR-10
 127 examples. Further details are in the appendix.

128 4.2 BN gain distinguishes the unstable scope settings

129 Figure 1 isolates the BN affine terms. At $\sigma = 1$, weights-only L2 retains 90.80% accuracy and
 130 adding BN- β retains 90.19%. Adding BN- γ instead falls to 22.50%; adding both affine terms
 131 reaches 24.41%. At $\sigma = 3$, the first two configurations still achieve 75.63% and 70.65%, while both
 132 configurations containing BN- γ are at chance. The near-overlap of the BN- γ and BN- $\gamma + \beta$ curves
 133 implicates gain, rather than bias, as the principal scope difference in this controlled run.

134 The scope result is consistent with the layerwise conversion diagnostics. In late VGG blocks, all-
 135 parameter L2 and weights-plus-BN- γ reach $\rho_l \approx 0.05$ – 0.11 and $P_{\text{cross},l} \approx 0.62$ – 0.74 , whereas MNE-
 136 L2 and weights-only L2 retain larger margins and mostly stay below 0.15 crossings (Appendix B).
 137 This does not establish a universal BN mechanism: the four-way intervention uses one training
 138 seed. It shows that reporting “L2” without its parameter scope can conceal experimentally important
 139 variation.

140 4.3 MNE-L2 across CIFAR and ImageNet

141 Figure 2(a,b) extends the stress test to five independently trained checkpoints per CIFAR dataset. At
 142 $\sigma = 3$, MNE-L2 reaches 86.59% on CIFAR-10 and 55.19% on CIFAR-100, compared with 82.01%
 143 and 45.33% for weights-only L1. At $\sigma = 5$, the gap widens to 72.06% versus 39.40% on CIFAR-10
 144 and 41.14% versus 17.91% on CIFAR-100. All-parameter L2 collapses near $\sigma = 1$ on both datasets,
 145 consistent with the scope ablation.

146 The high-noise ordering persists for VGG13 and VGG19 across both CIFAR datasets (Appendix C.2),
 147 indicating that it is not specific to VGG16 depth.

148 The single-checkpoint ImageNet result in Figure 2(c) follows the CIFAR ordering at moderate noise:
 149 MNE-L2 retains 33.07% at $\sigma = 1$, versus 21.11% for weights-only L1, and remains highest at $\sigma = 2$.
 150 This residual, larger-dataset extension is preliminary because it uses one seed.




Figure 2: Post-input-IF Gaussian-noise sweeps at $L = T = 16$. (a,b) CIFAR VGG16 curves are five-seed means; bands show sample standard deviations across seeds 40–44. (c) ImageNet ResNet-18 uses one seed-42 checkpoint and therefore has no uncertainty band. MNE-L2 is highlighted in red and all data points use circular markers.

Table 1: Five-seed CIFAR VGG16 robustness–activity summary at $T = 16$ (mean $\pm$ sample standard deviation; 10,000 test images per seed). Accuracy is measured under post-input-IF Gaussian noise at $\sigma = 5$; firing density and arithmetic energy use clean inference. The best value in each dataset and metric is bold, and the second-best is underlined. Clean accuracy and full activity distributions appear in Appendix C.1.

Dataset	Method	Clean firing density	Clean energy (mJ)	Acc. at $\sigma = 5$ (%)
CIFAR-10	L2-all	0.105 $\pm$ 0.003	0.585 $\pm$ 0.006	10.00 $\pm$ 0.00
	L2-wo	0.102 $\pm$ 0.001	0.505 $\pm$ 0.003	19.92 $\pm$ 4.29
	L1-wo	0.077 $\pm$ 0.001	0.402 $\pm$ 0.005	39.40 $\pm$ 2.60
	MNE-L2 (ours)	<u>0.083 $\pm$ 0.001</u>	<u>0.472 $\pm$ 0.003</u>	72.06 $\pm$ 2.57
CIFAR-100	L2-all	0.106 $\pm$ 0.003	0.558 $\pm$ 0.008	1.05 $\pm$ 0.06
	L2-wo	0.105 $\pm$ 0.001	0.522 $\pm$ 0.002	4.21 $\pm$ 0.24
	L1-wo	0.084 $\pm$ 0.002	0.429 $\pm$ 0.005	17.91 $\pm$ 4.84
	MNE-L2 (ours)	<u>0.090 $\pm$ 0.001</u>	<u>0.494 $\pm$ 0.004</u>	41.14 $\pm$ 1.23

4.4 Clean sparsity, robustness, and estimated energy

Table 1 compares all four regularizers at high noise. On CIFAR-10, MNE-L2 lowers clean firing from 0.105 to 0.083 and energy from 0.585 to 0.472 mJ relative to L2-all, while retaining 72.06% rather than 10.00% accuracy at $\sigma = 5$. The same robustness ordering holds on CIFAR-100, where MNE-L2 reaches 41.14% and L2-all reaches 1.05%. L1-wo is the sparsest and least energy-intensive method, but Appendix C.1 shows that it trails MNE-L2 by 6.49 clean CIFAR-100 points.

Following Horowitz [26], we charge repeated first-layer MACs at 4.6 pJ and measured event-triggered ACs in later layers at 0.9 pJ (Appendix Eq. 9). Only clean activity enters energy: CIFAR-10 L2-all firing at $T = 4$ rises from 0.105 to 0.139 under $\sigma = 1$, so noisy activity would confound injected events with deployment cost. Appendix C.1 shows that L2-wo is robust but nearly as active as L2-all; L1 is sparsest, but trails MNE-L2 by 6.49 clean CIFAR-100 points. These estimates omit memory, neuron updates, control, interconnect, and leakage, and are not device measurements.

5 Implications, Limitations, and Conclusion

For deployment, regularizer names should specify their parameter scope: including BN- γ changes the observed robustness ranking. MNE-L2 instead weights convolutional and linear parameters using folded BN statistics without shrinking affine or threshold terms, yielding lower activity than L2-all and an intermediate point between sparse but inaccurate L1 and active, clean-accurate L2-wo. Limitations include a one-seed scope study, one ImageNet checkpoint, synthetic single-site noise, unmatched clean accuracy, and arithmetic-only energy; five-seed CIFAR evidence remains restricted to VGG-family architectures. These constraints preclude universal dominance, but within the tested settings MNE-L2 better preserves post-IF representations than L2-all.

References

- [1] Mike Davies, Narayan Srinivasa, Tsung-Han Lin, Gautham Chinya, Yongqiang Cao, Sri Harsha Choday, Georgios Dimou, Prasad Joshi, Nabil Imam, Shweta Jain, et al. Loihi: A neuromorphic manycore processor with on-chip learning. *IEEE Micro*, 38(1):82–99, 2018. doi: 10.1109/MM.2018.112130359.
- [2] Bodo Rueckauer, Iulia-Alexandra Lungu, Yuhuang Hu, Michael Pfeiffer, and Shih-Chii Liu. Conversion of continuous-valued deep networks to efficient event-driven networks for image classification. *Frontiers in Neuroscience*, 11:682, 2017. doi: 10.3389/fnins.2017.00682.
- [3] Abhronil Sengupta, Yuting Ye, Robert Wang, Chiao Liu, and Kaushik Roy. Going deeper in spiking neural networks: VGG and residual architectures. *Frontiers in Neuroscience*, 13:95, 2019. doi: 10.3389/fnins.2019.00095.
- [4] Yangfan Hu, Qian Zheng, Xudong Jiang, and Gang Pan. Fast-SNN: Fast spiking neural network by converting quantized ANN. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 45(12):14546–14562, 2023.
- [5] Tong Bu, Wei Fang, Jianhao Ding, PengLin Dai, Zhaofei Yu, and Tiejun Huang. Optimal ANN-SNN conversion for high-accuracy and ultra-low-latency spiking neural networks. In *International Conference on Learning Representations*, 2022. URL https://openreview.net/forum?id=7B3IJMM1k_M.
- [6] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *International Conference on Machine Learning*, pages 448–456, 2015.
- [7] Sander M. Bohte, Joost N. Kok, and Han La Poutre. Error-backpropagation in temporally encoded networks of spiking neurons. *Neurocomputing*, 48(1–4):17–37, 2002.
- [8] Wei Fang, Zhaofei Yu, Yanqi Chen, Tiejun Huang, Timothée Masquelier, and Yonghong Tian. Deep residual learning in spiking neural networks. In *Advances in Neural Information Processing Systems*, volume 34, pages 21056–21069, 2021.
- [9] Yufei Guo, Xuhui Huang, and Zhe Ma. Direct learning-based deep spiking neural networks: A review. *Frontiers in Neuroscience*, 17:1209795, 2023.
- [10] Peter U. Diehl, Daniel Neil, Jonathan Binas, Matthew Cook, Shih-Chii Liu, and Michael Pfeiffer. Fast-classifying, high-accuracy spiking deep networks through weight and threshold balancing. In *International Joint Conference on Neural Networks*, pages 1–8, 2015.
- [11] Tong Bu, Jianhao Ding, Zhaofei Yu, and Tiejun Huang. Optimized potential initialization for low-latency spiking neural networks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pages 11–20, 2022.
- [12] Yuhang Li, Shikuang Deng, Xin Dong, Ruihao Gong, and Shi Gu. A free lunch from ANN: Towards efficient, accurate spiking neural networks calibration. In *International Conference on Machine Learning*, pages 6316–6325, 2021.
- [13] Zecheng Hao, Tong Bu, Jianhao Ding, Tiejun Huang, and Zhaofei Yu. Reducing ANN-SNN conversion error through residual membrane potential. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pages 11–21, 2023.
- [14] Zecheng Hao, Jianhao Ding, Tong Bu, Tiejun Huang, and Zhaofei Yu. Bridging the gap between ANNs and SNNs by calibrating offset spikes. In *International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=PFbz0WZyZRX>.
- [15] Velibor Bojkovic, Xiaofeng Wu, and Bin Gu. Temporal misalignment in ANN-SNN conversion and its mitigation via probabilistic spiking neurons. In *International Conference on Machine Learning*, 2025. URL <https://openreview.net/forum?id=Kip4avTjth>.

- [16] Fangxin Liu, Wenbo Zhao, Yongbiao Chen, Zongwu Wang, and Li Jiang. SpikeConverter: An efficient conversion framework zipping the gap between artificial neural networks and spiking neural networks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pages 1692–1701, 2022.
- [17] Chang Liu, Jiangrong Shen, Xuming Ran, Mingkun Xu, Qi Xu, Yi Xu, and Gang Pan. Efficient ANN-SNN conversion with error compensation learning. In *International Conference on Machine Learning*, 2025.
- [18] Yuchen Yang, Jingcheng Liu, Chengting Yu, Chengyi Yang, Gaoang Wang, and Aili Wang. NeuBridge: Bridging quantized activations and spiking neurons for ANN-SNN conversion. *Neuromorphic Computing and Engineering*, 5(2):024018, 2025.
- [19] Arash Rahnema, Andre T. Nguyen, and Edward Raff. Robust design of deep neural networks against adversarial attacks based on lyapunov theory. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 8178–8187, 2020.
- [20] Jianhao Ding, Zhiyu Pan, Yujia Liu, Zhaofer Yu, and Tiejun Huang. Robust stable spiking neural networks. In *Proceedings of the 41st International Conference on Machine Learning*, pages 11016–11029, 2024.
- [21] Ernesto Araya, Massimiliano Datres, and Gitta Kutyniok. Random spiking neural networks are stable and spectrally simple. In *International Conference on Learning Representations*, 2026.
- [22] Ozan Özdenizci and Robert Legenstein. Adversarially robust spiking neural networks through conversion. *Transactions on Machine Learning Research*, 2024. URL <https://openreview.net/forum?id=I8FMYa2BdP>.
- [23] Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- [24] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. ImageNet: A large-scale hierarchical image database. In *IEEE Conference on Computer Vision and Pattern Recognition*, pages 248–255, 2009.
- [25] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 770–778, 2016.
- [26] Mark Horowitz. 1.1 computing’s energy problem (and what we can do about it). In *IEEE International Solid-State Circuits Conference Digest of Technical Papers*, pages 10–14, 2014. doi: 10.1109/ISSCC.2014.6757323.
- [27] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998. doi: 10.1109/5.726791.
- [28] Han Xiao, Kashif Rasul, and Roland Vollgraf. Fashion-MNIST: A novel image dataset for benchmarking machine learning algorithms. *arXiv preprint arXiv:1708.07747*, 2017.

A A Local Precision–Width–Time Model

Consider an idealized normalized aggregation of independent perturbations,

$$\delta_{C,T} = \frac{1}{CT} \sum_{t=1}^T \sum_{i=1}^C w_i \epsilon_{i,t}, \quad \epsilon_{i,t} \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma^2). \quad (6)$$

If $C^{-1} \sum_i w_i^2 \leq M$, then

$$\text{Var}(\delta_{C,T}) = \frac{\sigma^2}{C^2 T} \sum_i w_i^2 \leq \frac{\sigma^2 M}{CT}. \quad (7)$$

Table 2: Five-seed clean MNIST conversion at $T = 4$ (mean $\pm$ standard deviation, no injected noise). Δ is $L=16$ minus $L=4$ SNN accuracy.

Channels	SNN, $L = 4$	SNN, $L = 16$	Δ
2/4	96.14 ± 0.91	87.80 ± 6.94	-8.34
4/8	98.00 ± 0.26	97.01 ± 0.60	-0.99
8/16	98.67 ± 0.13	98.22 ± 0.37	-0.44
16/32	98.67 ± 0.12	98.42 ± 0.21	-0.25




Figure 3: Layerwise CIFAR-10 VGG16 diagnostics at fixed $\sigma = 1$. All-parameter L2 and weights+BN- γ checkpoints develop small margin-to-noise ratios and high bin-crossing probability in later blocks.

For an unclipped QCFS cell, the nominal half-cell margin is $h_L = \lambda/(2L)$. Combining this value with the variance bound gives the nominal proxy

$$\rho_{\text{nom}} = \frac{h_L}{\sigma_{\text{eff}}} \geq \frac{\lambda\sqrt{CT}}{2L\sigma\sqrt{M}}. \quad (8)$$

Equation (2) follows because a changed quantization code requires $|\delta_l| \geq d_l(u)$ and a zero-mean Gaussian satisfies $\Pr(|\delta_l| \geq d) = 2Q(d/\sigma_{\text{eff},l})$. This local model isolates the competing directions of C , T , and L ; it is not a claim that all network perturbations are independent or Gaussian, nor a proof that the empirical accuracy gap must be monotone.

A.1 Compact precision–width control

To separate conversion error from injected-noise calibration, we train weight-decayed MNIST CNNs with channel pairs (2, 4), (4, 8), (8, 16), and (16, 32) at $L \in \{4, 16\}$. Each setting uses five training seeds and is converted at $T = 4$ with no injected noise. Raising L from 4 to 16 improves every source ANN by 0.06–0.29 points, yet reduces $T = 4$ SNN accuracy by 8.34 points for the 2/4-channel model. This loss contracts to 0.25 points at 16/32 channels (Table 2).

B Layerwise Diagnostics

Figure 3 reports the two diagnostics from Eq. (3). Layer 0 is the perturbed input IF; its effective-noise denominator is zero by construction for the downstream-propagation statistic and is omitted from the left panel. The right panel directly measures whether clean and noisy IF outputs differ.

C Additional Experimental Details

The CIFAR VGG16 source ANNs are trained for 300 epochs with batch size 128 and seeds 40–44 where multi-seed results are reported. Training uses SGD with momentum 0.9, initial learning rate 0.1, and cosine annealing. Random crops with four-pixel padding, horizontal flips, CIFAR AutoAugment,

Table 3: Top-1 accuracy under post-input-IF noise at $\sigma = 1$ (mean $\pm$ sample standard deviation across seeds 40–44).

Dataset	Method	$T = 4$	$T = 8$	$T = 16$
CIFAR-10	L2-all	10.10 ± 0.15	11.17 ± 1.18	23.30 ± 6.46
	MNE-L2 (ours)	88.97 ± 0.21	89.76 ± 0.18	90.09 ± 0.10
CIFAR-100	L2-all	1.56 ± 0.40	2.88 ± 0.83	14.64 ± 3.16
	MNE-L2 (ours)	58.77 ± 0.28	59.27 ± 0.11	59.48 ± 0.23

Table 4: Full clean SNN top-1 accuracy (mean $\pm$ sample standard deviation across seeds 40–44).

Dataset	Method	$T = 4$	$T = 8$	$T = 16$
CIFAR-10	L2-all	90.29 ± 0.23	90.52 ± 0.14	90.59 ± 0.13
	L2-wo	90.57 ± 0.18	90.96 ± 0.06	91.10 ± 0.07
	L1-wo	89.09 ± 0.25	89.70 ± 0.30	89.87 ± 0.20
	MNE-L2 (ours)	89.93 ± 0.16	90.31 ± 0.17	90.38 ± 0.14
CIFAR-100	L2-all	61.83 ± 0.47	62.22 ± 0.49	62.27 ± 0.37
	L2-wo	62.68 ± 0.54	63.57 ± 0.31	63.83 ± 0.26
	L1-wo	52.37 ± 0.95	52.89 ± 0.71	53.13 ± 0.73
	MNE-L2 (ours)	59.24 ± 0.21	59.62 ± 0.31	59.62 ± 0.22

and Cutout are used during training; evaluation uses only dataset-specific normalization. MNE-L2 uses coefficient 10^{-4} with detached BN affine/statistical factors and detached IF thresholds. All-parameter L2 uses weight decay 5×10^{-4} , weights-only L1 uses 10^{-5} , and the selective L2 scope controls use coefficient 2.5×10^{-4} . The scope comparison reuses CIFAR-10 seed-42 checkpoints and the same evaluation protocol at every severity.

The ImageNet ResNet-18 source ANNs are trained at $T = 0$ for 90 epochs with seed 42, batch size 256, SGD momentum 0.9, initial learning rate 0.1, and cosine annealing. Training applies a random 224×224 resized crop and horizontal flip; validation resizes to 256 pixels and takes a 224×224 center crop. Both splits use standard ImageNet normalization. The highest clean validation-accuracy checkpoint is converted with $L = T = 16$ and evaluated on the 50,000-image validation split. All-parameter and weights-only L2 use coefficient 10^{-4} , MNE-L2 uses 10^{-4} , and weights-only L1 uses 10^{-5} . No coefficient sweep or clean-accuracy matching was performed for the reported ImageNet comparison, and coefficients were not selected from the noisy evaluation curves.

For the CIFAR arithmetic estimate, dense ANN energy and event-driven SNN energy are

$$E_{\text{ANN}} = N_{\text{MAC}} E_{\text{MAC}}, \quad E_{\text{SNN}} = T N_{\text{MAC}}^{(1)} E_{\text{MAC}} + N_{\text{AC}}^{(>1)} E_{\text{AC}}, \quad (9)$$

with $E_{\text{MAC}} = 4.6$ pJ and $E_{\text{AC}} = 0.9$ pJ. The first convolution receives analog pixels at every timestep; all later convolutional and linear layers use measured presynaptic events. Only clean ($\sigma = 0$) activity is used in this estimate.

C.1 Full CIFAR robustness and energy statistics

Table 3 reports the full $\sigma = 1$ timestep sweep for the main comparison. Tables 4 and 5 report clean accuracy and energy at every timestep; Table 6 adds per-sample activity distributions and the L2-wo/L1 controls at $T = 16$. The main-paper Table 1 reports the four-method $\sigma = 5$ comparison. For energy accounting, we count dense operations in the first convolution at every timestep and event-triggered ACs in all later convolutional and linear layers. Every table uses all 10,000 examples in each test set; uncertainty reflects variation across five independently trained checkpoints, not measurement noise from physical hardware.

Table 5: Full clean CIFAR VGG16 arithmetic-energy estimates (mean $\pm$ sample standard deviation across seeds 40–44). ANN energy is 1.5277 mJ on CIFAR-10 and 1.5294 mJ on CIFAR-100. “Ratio” is ANN energy divided by estimated SNN energy.

Dataset	Method	$T = 4$		$T = 8$		$T = 16$	
		Energy (mJ)	Ratio	Energy (mJ)	Ratio	Energy (mJ)	Ratio
CIFAR-10	L2-all	0.1421 $\pm$ 0.0015	10.75 $\pm$ 0.11	0.2877 $\pm$ 0.0030	5.31 $\pm$ 0.05	0.5848 $\pm$ 0.0060	2.61 $\pm$ 0.03
	L2-wo	0.1227 $\pm$ 0.0008	12.45 $\pm$ 0.08	0.2483 $\pm$ 0.0016	6.15 $\pm$ 0.04	0.5050 $\pm$ 0.0033	3.03 $\pm$ 0.02
	L1-wo	0.0975 $\pm$ 0.0010	15.67 $\pm$ 0.17	0.1969 $\pm$ 0.0022	7.76 $\pm$ 0.08	0.4019 $\pm$ 0.0046	3.80 $\pm$ 0.04
	MNE-L2 (ours)	0.1141 $\pm$ 0.0005	13.39 $\pm$ 0.06	0.2312 $\pm$ 0.0012	6.61 $\pm$ 0.03	0.4716 $\pm$ 0.0026	3.24 $\pm$ 0.02
CIFAR-100	L2-all	0.1354 $\pm$ 0.0020	11.30 $\pm$ 0.16	0.2742 $\pm$ 0.0039	5.58 $\pm$ 0.08	0.5577 $\pm$ 0.0078	2.74 $\pm$ 0.04
	L2-wo	0.1271 $\pm$ 0.0005	12.03 $\pm$ 0.04	0.2569 $\pm$ 0.0010	5.95 $\pm$ 0.02	0.5225 $\pm$ 0.0021	2.93 $\pm$ 0.01
	L1-wo	0.1043 $\pm$ 0.0011	14.66 $\pm$ 0.16	0.2109 $\pm$ 0.0025	7.25 $\pm$ 0.09	0.4294 $\pm$ 0.0052	3.56 $\pm$ 0.04
	MNE-L2 (ours)	0.1197 $\pm$ 0.0009	12.77 $\pm$ 0.10	0.2424 $\pm$ 0.0019	6.31 $\pm$ 0.05	0.4942 $\pm$ 0.0037	3.09 $\pm$ 0.02

Table 6: Clean activity distribution, clean arithmetic energy, and $\sigma = 1$ robustness at $T = 16$. Within each seed, firing density is the fraction of nonzero entries across all IF outputs and is summarized over 10,000 samples by its mean, median, and 95th percentile; entries report the mean $\pm$ sample standard deviation of these statistics across seeds 40–44. Energy uses clean mean activity only.

Dataset	Method	Clean acc. (%)	Fire mean	Fire median	Fire p95	Energy (mJ)	Acc. at $\sigma = 1$ (%)
CIFAR-10	L2-all	90.59 $\pm$ 0.13	0.105 $\pm$ 0.003	0.105 $\pm$ 0.003	0.120 $\pm$ 0.004	0.585 $\pm$ 0.006	23.30 $\pm$ 6.46
	MNE-L2 (ours)	90.38 $\pm$ 0.14	0.083 $\pm$ 0.001	0.083 $\pm$ 0.001	0.099 $\pm$ 0.002	0.472 $\pm$ 0.003	90.09 $\pm$ 0.10
	L2-wo	91.10 $\pm$ 0.07	0.102 $\pm$ 0.001	0.103 $\pm$ 0.001	0.118 $\pm$ 0.001	0.505 $\pm$ 0.003	90.82 $\pm$ 0.11
	L1-wo	89.87 $\pm$ 0.20	0.077 $\pm$ 0.001	0.077 $\pm$ 0.001	0.090 $\pm$ 0.002	0.402 $\pm$ 0.005	89.60 $\pm$ 0.16
CIFAR-100	L2-all	62.27 $\pm$ 0.37	0.106 $\pm$ 0.003	0.106 $\pm$ 0.003	0.121 $\pm$ 0.004	0.558 $\pm$ 0.008	14.64 $\pm$ 3.16
	MNE-L2 (ours)	59.62 $\pm$ 0.22	0.090 $\pm$ 0.001	0.090 $\pm$ 0.001	0.107 $\pm$ 0.002	0.494 $\pm$ 0.004	59.48 $\pm$ 0.23
	L2-wo	63.83 $\pm$ 0.26	0.105 $\pm$ 0.001	0.106 $\pm$ 0.001	0.120 $\pm$ 0.001	0.522 $\pm$ 0.002	63.79 $\pm$ 0.25
	L1-wo	53.13 $\pm$ 0.73	0.084 $\pm$ 0.002	0.085 $\pm$ 0.002	0.101 $\pm$ 0.002	0.429 $\pm$ 0.005	53.08 $\pm$ 0.76

303 C.2 Depth transfer across VGG13 and VGG19

304 We repeat the post-input-IF Gaussian-noise sweep at $L = T = 16$ using VGG13 and VGG19 on
305 CIFAR-10 and CIFAR-100. Each curve reports the mean and sample standard deviation across
306 independently trained seeds 40–44. In addition to the four main regularizers, this control includes
307 Calibrated MNE with the fixed $\alpha = 0.1$ configuration; α is a prescribed training schedule rather than
308 a gradient-triggered or adaptive choice.

309 Across VGG13/16/19, MNE-L2 attains 74.28/72.06/73.23% at $\sigma = 5$ on CIFAR-10 and
310 44.01/41.14/37.60% on CIFAR-100. Calibrated MNE closely follows L2-wo despite matching
311 or slightly exceeding its clean accuracy, while L1-wo incurs a substantial CIFAR-100 clean-accuracy
312 cost that grows from VGG13 to VGG19. Thus the high-noise MNE-L2 advantage transfers across
313 the three completed depths, although the CIFAR-100 endpoint declines with depth.




Figure 4: Depth-transfer robustness at $L = T = 16$. Gaussian noise is injected after the input IF. Curves are five-seed means and bands show sample standard deviations. MNE-L2 is highlighted in red; Calibrated MNE is included only as an appendix control.

Table 7: Additional-depth accuracy endpoints (mean $\pm$ sample standard deviation across seeds 40–44). The best value within each architecture, dataset, and noise condition is bold.

Architecture	Method	CIFAR-10		CIFAR-100	
		$\sigma = 0$	$\sigma = 5$	$\sigma = 0$	$\sigma = 5$
VGG13	L2-all	90.93 $\pm$ 0.14	9.99 $\pm$ 0.01	64.73 $\pm$ 0.19	1.01 $\pm$ 0.01
	L2-wo	91.44 $\pm$ 0.14	22.78 $\pm$ 5.99	65.94 $\pm$ 0.15	5.96 $\pm$ 0.96
	L1-wo	90.30 $\pm$ 0.14	42.17 $\pm$ 10.20	51.16 $\pm$ 1.31	32.06 $\pm$ 3.92
	MNE-L2 (ours)	90.61 $\pm$ 0.19	74.28 $\pm$ 3.25	62.41 $\pm$ 0.15	44.01 $\pm$ 2.02
	Cal. MNE	91.41 $\pm$ 0.13	21.38 $\pm$ 2.58	65.82 $\pm$ 0.22	5.59 $\pm$ 0.92
VGG19	L2-all	90.19 $\pm$ 0.12	10.00 $\pm$ 0.00	58.37 $\pm$ 0.64	1.13 $\pm$ 0.18
	L2-wo	90.69 $\pm$ 0.18	19.95 $\pm$ 4.44	60.62 $\pm$ 0.32	4.03 $\pm$ 0.65
	L1-wo	88.09 $\pm$ 0.25	53.04 $\pm$ 3.85	45.04 $\pm$ 1.01	9.77 $\pm$ 5.59
	MNE-L2 (ours)	89.67 $\pm$ 0.06	73.23 $\pm$ 2.50	55.75 $\pm$ 0.43	37.60 $\pm$ 3.19
	Cal. MNE	90.84 $\pm$ 0.03	23.66 $\pm$ 3.05	60.73 $\pm$ 0.22	4.73 $\pm$ 0.76

314 C.3 Compact MNIST-family controls

315 We use MNIST [27] and Fashion-MNIST [28] with a 2,081-parameter CNN: two 3×3 convolution-
 316 BN-IF blocks with 2 and 4 channels, 2×2 max pooling, and a linear classifier. Source ANNs are
 317 trained for 30 epochs with seeds 40–44 and converted at $T \in \{4, 8, 16\}$. We compare MNE-L2,
 318 orthogonal weight regularization, no regularization, and optimizer L2. During the low-latency sweep,
 319 Gaussian noise is added after both the input IF and first hidden IF at $\sigma \in \{0, 0.25, 0.5\}$; each nonzero
 320 point averages three paired streams within each training seed.

321 For the compact CNN, the dense ANN baseline is 30,184 MACs per sample. The event counter
 322 charges every nonzero presynaptic activation once for each valid downstream connection. It counts
 323 both convolutional layers and the classifier, including edge effects from padding, and accumulates
 324 over all T timesteps. Reported means and standard deviations are across five independently trained
 325 checkpoints.

326 Figure 5 gives the compact low-latency view. At $T = 4$, the event-SynOps ratio ranges from 0.36 to
 327 0.45 on Fashion-MNIST and 0.63 to 0.65 on MNIST; at $T = 16$, it rises to 1.54–1.97 and 2.71–2.86.
 328 These are algorithmic operation counts rather than wall-clock latency or hardware energy [1, 26].




Figure 5: Appendix compact-model control. Top: five-seed accuracy under two-site Gaussian noise with $\sigma = 0.5$. Bottom: event-driven synaptic-operation proxy relative to one dense ANN MAC pass; the dashed line marks ratio 1. Error bars are standard deviations across training seeds.

Table 8: Full five-seed accuracy at $\sigma = 0.5$ (mean $\pm$ standard deviation across training seeds). Noise is injected after both the input IF and first hidden IF.

Dataset	Method	$T = 4$	$T = 8$	$T = 16$
Fashion	MNE-L2 (ours)	42.01 ± 6.48	50.81 ± 11.93	55.19 ± 12.53
Fashion	Orthogonal	41.98 ± 5.49	50.50 ± 9.54	55.12 ± 10.95
Fashion	No Reg	37.53 ± 4.69	45.53 ± 9.30	51.64 ± 11.19
Fashion	L2	30.75 ± 7.49	34.35 ± 9.95	38.06 ± 11.80
MNIST	MNE-L2 (ours)	74.55 ± 16.12	73.44 ± 17.98	67.54 ± 18.60
MNIST	Orthogonal	74.62 ± 17.89	74.99 ± 20.82	70.17 ± 22.92
MNIST	No Reg	69.75 ± 13.68	69.44 ± 15.69	65.31 ± 16.18
MNIST	L2	29.41 ± 19.27	30.00 ± 19.23	30.30 ± 19.07

D Reproducibility Notes

All reported datasets use their standard public train/test splits. The evaluation code fixes the training seed separately from the perturbation seed, stores every per-seed result before aggregation, and reports sample standard deviation rather than standard error. Figures 1, 2, and 4 are generated directly from stored CSV summaries by a deterministic plotting script; curve values are neither redrawn nor transcribed. The compact event-count control uses CPU inference and 1,024 examples per checkpoint. The CIFAR activity evaluation uses each full 10,000-example test set: sample mean, median, and p95 are computed within each seed, then summarized across seeds 40–44 by mean and sample standard deviation. Only $\sigma = 0$ activity enters energy; $\sigma = 1$ activity is retained as a diagnostic and excluded from deployment estimates. Full VGG and ResNet training requires GPU compute; evaluation does not update model parameters.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract, Introduction, Experiments, and Limitations state the MNE-L2 contribution, the single-seed BN-scope ablation, the five-seed CIFAR evidence, the single-checkpoint ImageNet extension, the clean-accuracy trade-off, and the non-hardware nature of the operation proxy.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The Implications and Limitations section identifies the single-seed BN-scope study, the single-checkpoint ImageNet extension, synthetic post-IF perturbations, restricted multi-seed evidence, clean-accuracy trade-off, and absence of hardware measurements.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Section 2 states the local Gaussian boundary-crossing bound and Appendix A derives the aggregation variance bound while explicitly listing the independence, Gaussianity, and bounded-weight assumptions and limiting the claim to the idealized local model.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The Experiments setup and appendix specify datasets, architectures, data transforms, optimizers, schedules, batch sizes, seeds, checkpoint selection, perturbation sites, aggregation, and the event-counting rule.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: The code and result tables exist, but an anonymized archival artifact with environment and end-to-end commands has not yet been attached to this draft.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The Experiments setup and appendix report the train/test splits, transforms, optimizers, schedules, batch sizes, epochs, coefficients, checkpoint selection, and evaluation protocol. They also state that the reported ImageNet coefficients were not clean-accuracy matched or selected from noisy evaluation curves.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Figures report sample standard deviations; the text distinguishes five training seeds from five paired perturbation streams and states the aggregation unit.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).

- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: The appendix identifies CPU versus GPU workloads, but exact accelerator type, runtime, and total training compute still need to be recovered from cluster logs.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work uses public benchmark data, evaluates technical reliability, preserves submission anonymity, and does not involve people or sensitive data.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The Limitations section discusses the benefit of more reliable reporting and the risk of overstating deployment robustness.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.

- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The work does not release a high-risk model, scraped dataset, or dual-use generative system.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [No]

Justification: The benchmark datasets and prior methods are cited, but this draft does not yet list every asset license and version explicitly.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [No]

Justification: New evaluation and plotting code is present in the working artifact, but release documentation and a license are not yet finalized.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The research does not use crowdsourcing or human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The research does not use crowdsourcing or human subjects and therefore does not require IRB review.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

- 659 • Depending on the country in which research is conducted, IRB approval (or equivalent)
660 may be required for any human subjects research. If you obtained IRB approval, you
661 should clearly state this in the paper.
- 662 • We recognize that the procedures for this may vary significantly between institutions
663 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
664 guidelines for their institution.
- 665 • For initial submissions, do not include any information that would break anonymity (if
666 applicable), such as the institution conducting the review.

667 16. **Declaration of LLM usage**

668 Question: Does the paper describe the usage of LLMs if it is an important, original, or
669 non-standard component of the core methods in this research? Note that if the LLM is used
670 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
671 scientific rigor, or originality of the research, declaration is not required.

672 Answer: [N/A]

673 Justification: LLMs are not an important, original, or non-standard component of the
674 research method.

675 Guidelines:

- 676 • The answer [N/A] means that the core method development in this research does not
677 involve LLMs as any important, original, or non-standard components.
- 678 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
679 be described.